

# DSC520 Week11-12 Exercise 11.2.1

Anjani Bonda

March 4th 2022

```
# Load the packages
library(caTools)
library(ggplot2)
setwd("/Users/anjanibonda/DSC520/dsc520")
# Load the binary classifier dataset to dataframe
binary_df <- read.csv("data/binary-classifier-data.csv")
# Examine the structure
str(binary_df)
```

```
## 'data.frame': 1498 obs. of 3 variables:
## $ label: int 0 0 0 0 0 0 0 0 0 0 ...
## $ x : num 70.9 75 73.8 66.4 69.1 ...
## $ y : num 83.2 87.9 92.2 81.1 84.5 ...
```

```
# Check sample rows
head(binary_df)
```

```
## label x y
## 1 0 70.88469 83.17702
## 2 0 74.97176 87.92922
## 3 0 73.78333 92.20325
## 4 0 66.40747 81.10617
## 5 0 69.07399 84.53739
## 6 0 72.23616 86.38403
```

```
# Load the trinary classifier dataset to dataframe
trinary_df <- read.csv("data/trinary-classifier-data.csv")
# Examine the structure
str(trinary_df)
```

```
## 'data.frame': 1568 obs. of 3 variables:
## $ label: int 0 0 0 0 0 0 0 0 0 0 ...
## $ x : num 30.1 31.3 34.1 32.6 34.7 ...
## $ y : num 39.6 51.8 49.3 41.2 45.5 ...
```

```
# Check sample rows
head(trinary_df)
```

```
## label x y
```

```
## 1    0 30.08387 39.63094
## 2    0 31.27613 51.77511
## 3    0 34.12138 49.27575
## 4    0 32.58222 41.23300
## 5    0 34.65069 45.47956
## 6    0 33.80513 44.24656
```


*# i. Plot the data from each dataset using scatter plot*

```
ggplot(binary_df, aes(x=binary_df$x, y=binary_df$y)) + geom_point(aes(color = binary_df$label))
```

```
## Warning: Use of 'binary_df$label' is discouraged. Use 'label' instead.
```

```
## Warning: Use of 'binary_df$x' is discouraged. Use 'x' instead.
```

```
## Warning: Use of 'binary_df$y' is discouraged. Use 'y' instead.
```



```
ggplot(trinary_df, aes(x=trinary_df$x, y=trinary_df$y)) + geom_point(aes(color = trinary_df$label))
```

```
## Warning: Use of 'trinary_df$label' is discouraged. Use 'label' instead.
```

```
## Warning: Use of 'trinary_df$x' is discouraged. Use 'x' instead.
```

```
## Warning: Use of 'trinary_df$y' is discouraged. Use 'y' instead.
```



```
# Normalization of binary_df
normalize <- function(x) { return ((x - min(x)) / (max(x) - min(x))) }
binary_df.n = as.data.frame(lapply(binary_df[,2:3], normalize))
binary_df.n = as.data.frame(lapply(binary_df[,2:3], normalize))
set.seed(123)
# Random sample of 70% of data
dat.d <- sample(1:nrow(binary_df.n), size = nrow(binary_df.n)*0.7, replace = FALSE)
# Create test and train datasets
train.binary_df <- binary_df[dat.d,]
test.binary_df <- binary_df[-dat.d,]
# Create separate dataframe for label feature
train.binary_df_label <- binary_df[dat.d,1]
test.binary_df_label <- binary_df[-dat.d,1]
# Find no. of observations
NROW(train.binary_df)
```

```
## [1] 1048
```

```
# From above, we have 700 observations in our training dataset. The square root of 700 is about 26.45.
# Therefore, we'll create 2 models. One with 'K' value 26 and other with 'K' value 27
library(class)
knn.binary_df.1 <- knn(train=train.binary_df, test=test.binary_df, cl=train.binary_df_label, k=1)
# Calculate accuracy of the models
# Caculate the proportion of correct classification for k=32,33
ACC.binary_df.1 <- 100*sum(test.binary_df_label == knn.binary_df.1)/NROW(test.binary_df_label)
ACC.binary_df.1
```

```
## [1] 98.22222
```

```
# Accuracy is 98.22  
# Check prediction against actual value in tabular form for k=32  
table(knn.binary_df.1, test.binary_df_label)
```

```
##               test.binary_df_label  
## knn.binary_df.1  0    1  
##               0 227    4  
##               1   4 215
```

```
# Use confusion matrix to calculate the accuracy.  
library(caret)
```

```
## Loading required package: lattice
```

```
confusionMatrix(table(knn.binary_df.1, test.binary_df_label))
```

```
## Confusion Matrix and Statistics  
##  
##               test.binary_df_label  
## knn.binary_df.1  0    1  
##               0 227    4  
##               1   4 215  
##  
##               Accuracy : 0.9822  
##               95% CI : (0.9653, 0.9923)  
##      No Information Rate : 0.5133  
##      P-Value [Acc > NIR] : <2e-16  
##  
##               Kappa : 0.9644  
##  
##      McNemar's Test P-Value : 1  
##  
##               Sensitivity : 0.9827  
##               Specificity : 0.9817  
##               Pos Pred Value : 0.9827  
##               Neg Pred Value : 0.9817  
##               Prevalence : 0.5133  
##               Detection Rate : 0.5044  
##      Detection Prevalence : 0.5133  
##               Balanced Accuracy : 0.9822  
##  
##               'Positive' Class : 0  
##
```

```
# Normalization of trinary_df  
normalize <- function(x) { return ((x - min(x)) / (max(x) - min(x))) }  
trinary_df.n = as.data.frame(lapply(trinary_df[,2:3], normalize))
```

```

trinary_df.n = as.data.frame(lapply(trinary_df[,2:3], normalize))
set.seed(123)
# Random sample of 70% of data
dat.d <- sample(1:nrow(trinary_df.n), size = nrow(trinary_df.n)*0.7, replace = FALSE)
# Create test and train datasets
train.trinary_df <- trinary_df[dat.d,]
test.trinary_df <- trinary_df[-dat.d,]
# Create separate dataframe for label feature
train.trinary_df_label <- trinary_df[dat.d,1]
test.trinary_df_label <- trinary_df[-dat.d,1]
# Find no. of observations
NROW(train.trinary_df)

```

```
## [1] 1097
```

```

library(class)
knn.trinary_df.1 <- knn(train=train.trinary_df, test=test.trinary_df, cl=train.trinary_df_label, k=1)
# Calculate accuracy of the models
# Caculate the proportion of correct classification for k=32,33
ACC.trinary_df.1 <- 100*sum(test.trinary_df_label == knn.trinary_df.1)/NROW(test.trinary_df_label)
ACC.trinary_df.1

```

```
## [1] 95.75372
```

```

# Accuracy is 95.75
# Check prediction against actual value in tabular form for k=32
table(knn.trinary_df.1, test.trinary_df_label)

```

```

##           test.trinary_df_label
## knn.trinary_df.1  0   1   2
##           0 131   7   1
##           1   3 185   7
##           2   0   2 135

```

```

# Use confusion matrix to calculate the accuracy.
library(caret)
confusionMatrix(table(knn.trinary_df.1, test.trinary_df_label))

```

```

## Confusion Matrix and Statistics
##
##           test.trinary_df_label
## knn.trinary_df.1  0   1   2
##           0 131   7   1
##           1   3 185   7
##           2   0   2 135
##
## Overall Statistics
##
##           Accuracy : 0.9575
##           95% CI : (0.9352, 0.9739)
##           No Information Rate : 0.4119

```

```
##      P-Value [Acc > NIR] : <2e-16
##
##              Kappa : 0.9354
##
## McNemar's Test P-Value : 0.1461
##
## Statistics by Class:
##
##              Class: 0 Class: 1 Class: 2
## Sensitivity      0.9776  0.9536  0.9441
## Specificity      0.9763  0.9639  0.9939
## Pos Pred Value   0.9424  0.9487  0.9854
## Neg Pred Value   0.9910  0.9674  0.9760
## Prevalence       0.2845  0.4119  0.3036
## Detection Rate   0.2781  0.3928  0.2866
## Detection Prevalence 0.2951  0.4140  0.2909
## Balanced Accuracy 0.9769  0.9588  0.9690
```

```
# ii Fit a k nearest neighbors' model for each dataset for k=3, k=5, k=10, k=15, k=20, and k=25. Compute
# Accuracy level of binary dataset
```


```
j <- 1
k.optm <- 1
for (i in c(3,5,10,15,20,25)){
  knn.mod <- knn(train=train.binary_df, test=test.binary_df, cl=train.binary_df_label, k=i)
  k.optm[i] <- 100 * sum(test.binary_df_label == knn.mod)/NROW(test.binary_df_label)
  k <- i
  j <- j + 1
  cat(k, '=', k.optm[i], ' ')}

```

```
## 3 = 98 5 = 97.55556 10 = 98.22222 15 = 97.55556 20 = 97.55556 25 = 98.44444
```

```
# Accuracy Plot
```

```
plot(k.optm, type="b", xlab="K-value", ylab="Accuracy of Binary Dataset")
```



```
# Accuracy level of trinary dataset
j <- 1
k.optm <- 1
for (i in c(3,5,10,15,20,25)){
  knn.mod <- knn(train=train.trinary_df, test=test.trinary_df, cl=train.trinary_df_label, k=i)
  k.optm[i] <- 100 * sum(test.trinary_df_label == knn.mod)/NROW(test.trinary_df_label)
  k <- i
  j <- j + 1
  cat(k, '=', k.optm[i], '\n')}


```

```
## 3 = 93.20594 5 = 92.14437 10 = 89.38429 15 = 89.17197 20 = 87.89809 25 = 86.83652

```

```
# Accuracy Plot
plot(k.optm, type="b", xlab="K-value", ylab="Accuracy of Trinary Dataset")

```



*# i. Looking back at the plots of the data, do you think a linear classifier would work well on these data?*


```
x1=binary_df[2]
x2=binary_df[3]
y <- sign(3 * x1 - 4 * x2 - 1)
y[y == -1] <- 0
df <- cbind.data.frame(y,x1,x2)
names(df)[1] <- 'y'
names(df)[2] <- 'x1'
names(df)[3] <- 'x2'
mdl <- glm(y ~ . , data=df,family = binomial)
```

```
## Warning: glm.fit: algorithm did not converge
```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```
slope <- coef(mdl)[2]/(-coef(mdl)[3])
intercept <- coef(mdl)[1]/(-coef(mdl)[3])
library(lattice)
xyplot(x2 ~ x1,data=df, groups=y,panel=function(...){
  panel.xyplot(...)
  panel.abline(intercept, slope)
  panel.grid(...)
})
```






```
x1= trinary_df[2]
x2= trinary_df[3]
y <- sign(3 * x1 - 4 * x2 - 1)
y[y == -1] <- 0
df <- cbind.data.frame(y,x1,x2)
names(df)[1] <- 'y'
names(df)[2] <- 'x1'
names(df)[3] <- 'x2'
mdl <- glm(y ~ . , data=df,family = binomial)
```

```
## Warning: glm.fit: algorithm did not converge
```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```
slope <- coef(mdl)[2]/(-coef(mdl)[3])
intercept <- coef(mdl)[1]/(-coef(mdl)[3])
library(lattice)
xyplot(x2 ~ x1,data=df, groups=y,panel=function(...){
  panel.xyplot(...)
  panel.abline(intercept, slope)
  panel.grid(...)
})
```



```
# Looking at the plots, I think that the linear classifier would work well on binary dataset but not on
# dataset
# ii. How does the accuracy of your logistic regression classifier from last week compare? Why is the

## The accuracy of logistic regression model was 67% but the accuracy of knn model is 98% of binary dat
## The difference in accuracy is due to the non-linearness of the data in the input datasets.
## KNN fits good for the non-linear dataset and hence it is more suitable model in our case.
```