

# Assignment 05

Author: Anjani Bonda

Date: 4/15/2023

## Assignment 5.1

Implement the movie review classifier found in section 3.4 of Deep Learning with Python.

```
In [27]: # Load all the required libraries
import numpy as np
from tensorflow import keras
from tensorflow.keras import layers
import matplotlib.pyplot as plt
```

```
In [28]: ## Import models and layers from keras
from keras import models
from keras import layers
from keras.utils import to_categorical
```

### Load the dataset

```
In [3]: # Load the IMDB dataset
from keras.datasets import imdb
# split train / test data
(train_data, train_labels), (test_data, test_labels) = imdb.load_data(
    num_words=10000)
```

Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/imdb.npz>  
17464789/17464789 [=====] - 1s 0us/step

### Prepare data

```
In [4]: # Encode the integer sequences via multi-hot encoding
def vectorize_sequences(sequences, dimension=10000):
    results = np.zeros((len(sequences), dimension))
    for i, sequence in enumerate(sequences):
        results[i, sequence] = 1.
    return results
```

```
In [5]: # Vectorize the data
x_train = vectorize_sequences(train_data)
x_test = vectorize_sequences(test_data)
```

```
In [6]: # Vectorize the labels
y_train = np.asarray(train_labels).astype("float32")
y_test = np.asarray(test_labels).astype("float32")
```

### Model Building

```
In [8]: ## Define keras model as sequential
model = models.Sequential()
model.add(layers.Dense(16, activation='relu', input_shape=(10000,)))
model.add(layers.Dense(16, activation='relu'))
model.add(layers.Dense(1, activation='sigmoid'))
```

```
In [9]: # Compile the model
model.compile(optimizer="rmsprop",
              loss="binary_crossentropy",
              metrics=["accuracy"])
```

## Model Validation

```
In [10]: ## Set aside a validation set (10,000 samples)
# Data
x_val = x_train[:10000]
partial_x_train = x_train[10000:]
# Labels
y_val = y_train[:10000]
partial_y_train = y_train[10000:]
```


## Model Training

```
In [11]: # Train the model
history = model.fit(partial_x_train,
                    partial_y_train,
                    epochs=20,
                    batch_size=512,
                    validation_data=(x_val, y_val))
```


```
Epoch 1/20
30/30 [=====] - 4s 110ms/step - loss: 0.5206 - accuracy: 0.7805 - val_loss: 0.4386 - val_accuracy: 0.8120
Epoch 2/20
30/30 [=====] - 1s 29ms/step - loss: 0.3208 - accuracy: 0.9011 - val_loss: 0.3271 - val_accuracy: 0.8760
Epoch 3/20
30/30 [=====] - 1s 32ms/step - loss: 0.2340 - accuracy: 0.9273 - val_loss: 0.2819 - val_accuracy: 0.8931
Epoch 4/20
30/30 [=====] - 1s 29ms/step - loss: 0.1840 - accuracy: 0.9394 - val_loss: 0.2781 - val_accuracy: 0.8896
Epoch 5/20
30/30 [=====] - 1s 27ms/step - loss: 0.1494 - accuracy: 0.9529 - val_loss: 0.2852 - val_accuracy: 0.8851
Epoch 6/20
30/30 [=====] - 1s 27ms/step - loss: 0.1233 - accuracy: 0.9626 - val_loss: 0.2891 - val_accuracy: 0.8882
Epoch 7/20
30/30 [=====] - 2s 54ms/step - loss: 0.1040 - accuracy: 0.9683 - val_loss: 0.3032 - val_accuracy: 0.8840
Epoch 8/20
30/30 [=====] - 1s 34ms/step - loss: 0.0852 - accuracy: 0.9762 - val_loss: 0.3599 - val_accuracy: 0.8724
Epoch 9/20
30/30 [=====] - 1s 41ms/step - loss: 0.0734 - accuracy: 0.9798 - val_loss: 0.3437 - val_accuracy: 0.8792
Epoch 10/20
30/30 [=====] - 1s 29ms/step - loss: 0.0622 - accuracy: 0.9831 - val_loss: 0.3723 - val_accuracy: 0.8786
Epoch 11/20
30/30 [=====] - 1s 33ms/step - loss: 0.0520 - accuracy: 0.9875 - val_loss: 0.4095 - val_accuracy: 0.8729
Epoch 12/20
30/30 [=====] - 1s 27ms/step - loss: 0.0417 - accuracy: 0.9911 - val_loss: 0.4306 - val_accuracy: 0.8711
Epoch 13/20
30/30 [=====] - 1s 30ms/step - loss: 0.0363 - accuracy: 0.9927 - val_loss: 0.4700 - val_accuracy: 0.8688
Epoch 14/20
30/30 [=====] - 1s 39ms/step - loss: 0.0306 - accuracy: 0.9935 - val_loss: 0.4750 - val_accuracy: 0.8732
Epoch 15/20
30/30 [=====] - 1s 28ms/step - loss: 0.0244 - accuracy: 0.9956 - val_loss: 0.5092 - val_accuracy: 0.8692
Epoch 16/20
30/30 [=====] - 1s 28ms/step - loss: 0.0182 - accuracy: 0.9976 - val_loss: 0.5584 - val_accuracy: 0.8646
Epoch 17/20
30/30 [=====] - 1s 27ms/step - loss: 0.0176 - accuracy: 0.9975 - val_loss: 0.5776 - val_accuracy: 0.8679
Epoch 18/20
30/30 [=====] - 1s 28ms/step - loss: 0.0128 - accuracy: 0.9987 - val_loss: 0.6228 - val_accuracy: 0.8643
Epoch 19/20
30/30 [=====] - 1s 27ms/step - loss: 0.0135 - accuracy: 0.9979 - val_loss: 0.6382 - val_accuracy: 0.8679
Epoch 20/20
30/30 [=====] - 1s 28ms/step - loss: 0.0062 - accuracy: 0.9997 - val_loss: 0.6653 - val_accuracy: 0.8661
```

## Plot Model Output and Loss

```
In [12]: # Plot the training and validation LOSS
history_dict = history.history
loss_values = history_dict["loss"]
val_loss_values = history_dict["val_loss"]
epochs = range(1, len(loss_values) + 1)
plt.plot(epochs, loss_values, "bo", label="Training loss") # 'bo' blue dot
plt.plot(epochs, val_loss_values, "b", label="Validation loss") # 'b' blue line
plt.title("Training and validation loss")
plt.xlabel("Epochs")
plt.ylabel("Loss")
plt.legend()
plt.show()
```



```
In [13]: # Plot the training and validation ACCURACY
plt.clf() # clear the figure
acc = history_dict["accuracy"]
val_acc = history_dict["val_accuracy"]
plt.plot(epochs, acc, "bo", label="Training acc")
plt.plot(epochs, val_acc, "b", label="Validation acc")
plt.title("Training and validation accuracy")
plt.xlabel("Epochs")
plt.ylabel("Accuracy")
plt.legend()
plt.show()
```



## Model Fitting

```
In [14]: # Fit model
model.fit(x_train, y_train, epochs=4, batch_size=512)

Epoch 1/4
49/49 [=====] - 2s 30ms/step - loss: 0.2287 - accuracy: 0.9454
Epoch 2/4
49/49 [=====] - 1s 18ms/step - loss: 0.1411 - accuracy: 0.9585
Epoch 3/4
49/49 [=====] - 1s 18ms/step - loss: 0.1156 - accuracy: 0.9653
Epoch 4/4
49/49 [=====] - 1s 18ms/step - loss: 0.0947 - accuracy: 0.9704

Out[14]: <keras.callbacks.History at 0x7fce6fa21f10>
```

```
In [15]: # Evaluate model
results = model.evaluate(x_test, y_test)

782/782 [=====] - 2s 3ms/step - loss: 0.4833 - accuracy: 0.8586
```

```
In [16]: # Show results
print(f'Test loss: {results[0]:0.3f}\nTest accuracy: {results[1]:0.3f}')
```

Test loss: 0.483  
Test accuracy: 0.859

## Generate predictions

```
In [17]: # Predict test data
model.predict(x_test)

782/782 [=====] - 3s 4ms/step
```

```
Out[17]: array([[0.01214925],
               [0.99996805],
               [0.71127635],
               ...,
               [0.03993864],
               [0.04501442],
               [0.7804223 ]], dtype=float32)
```

## Assignment 5.2

Implement the news classifier found in section 3.5 of Deep Learning with Python.

### Load the data

```
In [51]: # Load the Reuters dataset
from tensorflow.keras.datasets import reuters
# split train / test data
(train_data, train_labels), (test_data, test_labels) = reuters.load_data(
    num_words=10000)
```

```
In [52]: print("The length of training dataset: {}".format(len(train_data)))
print("The length of test dataset: {}".format(len(test_data)))
```

The length of training dataset: 8982

The length of test dataset: 2246

### Data Preparation

```
In [53]: # Encode the integer sequences via multi-hot encoding
def vectorize_sequences(sequences, dimension=10000):
    results = np.zeros((len(sequences), dimension))
    for i, sequence in enumerate(sequences):
        results[i, sequence] = 1.
    return results
```

```
In [54]: # Vectorize the input data
x_train = vectorize_sequences(train_data)
x_test = vectorize_sequences(test_data)
```

```
In [55]: # Vectorize the labels
y_train = np.asarray(train_labels).astype("float32")
y_test = np.asarray(test_labels).astype("float32")
```

```
In [56]: # Set up one-hot encoding
def to_one_hot(labels, dimension=46):
    results = np.zeros((len(labels), dimension))
    for i, label in enumerate(labels):
        results[i, label] = 1.
    return results
```

```
In [57]: # One-hot encode the labels
one_hot_train_labels = to_one_hot(train_labels)
one_hot_test_labels = to_one_hot(test_labels)
```

### Model Building

```
In [58]: # Define the model
model = models.Sequential()
model.add(layers.Dense(64, activation='relu', input_shape=(10000,)))
model.add(layers.Dense(64, activation='relu'))
model.add(layers.Dense(46, activation='softmax'))
```

```
In [59]: # Compile the model
model.compile(optimizer="rmsprop",
              loss="categorical_crossentropy",
              metrics=["accuracy"])
```

## Model Validation

```
In [60]: ## Set aside a validation set (1000 samples)
# Data
x_val = x_train[:1000]
partial_x_train = x_train[1000:]
# Labels
y_val = one_hot_train_labels[:1000]
partial_y_train = one_hot_train_labels[1000:]
```


```
In [61]: # Train the model
history = model.fit(partial_x_train,
                    partial_y_train,
                    epochs=20,
                    batch_size=512,
                    validation_data=(x_val, y_val))
```

```
Epoch 1/20
16/16 [=====] - 2s 58ms/step - loss: 2.5454 - accurac
y: 0.5256 - val_loss: 1.6854 - val_accuracy: 0.6400
Epoch 2/20
16/16 [=====] - 1s 40ms/step - loss: 1.3972 - accurac
y: 0.7050 - val_loss: 1.3175 - val_accuracy: 0.7050
Epoch 3/20
16/16 [=====] - 1s 41ms/step - loss: 1.0585 - accurac
y: 0.7720 - val_loss: 1.1556 - val_accuracy: 0.7520
Epoch 4/20
16/16 [=====] - 1s 48ms/step - loss: 0.8413 - accurac
y: 0.8239 - val_loss: 1.0654 - val_accuracy: 0.7670
Epoch 5/20
16/16 [=====] - 1s 56ms/step - loss: 0.6772 - accurac
y: 0.8502 - val_loss: 0.9840 - val_accuracy: 0.7890
Epoch 6/20
16/16 [=====] - 1s 47ms/step - loss: 0.5454 - accurac
y: 0.8806 - val_loss: 0.9450 - val_accuracy: 0.8040
Epoch 7/20
16/16 [=====] - 1s 57ms/step - loss: 0.4404 - accurac
y: 0.9067 - val_loss: 0.9305 - val_accuracy: 0.7970
Epoch 8/20
16/16 [=====] - 1s 51ms/step - loss: 0.3559 - accurac
y: 0.9222 - val_loss: 0.9143 - val_accuracy: 0.7990
Epoch 9/20
16/16 [=====] - 1s 47ms/step - loss: 0.2952 - accurac
y: 0.9362 - val_loss: 0.9145 - val_accuracy: 0.8080
Epoch 10/20
16/16 [=====] - 1s 65ms/step - loss: 0.2463 - accurac
y: 0.9442 - val_loss: 0.9470 - val_accuracy: 0.8100
Epoch 11/20
16/16 [=====] - 1s 76ms/step - loss: 0.2120 - accurac
y: 0.9481 - val_loss: 0.9420 - val_accuracy: 0.8010
Epoch 12/20
16/16 [=====] - 1s 54ms/step - loss: 0.1840 - accurac
y: 0.9514 - val_loss: 0.9235 - val_accuracy: 0.8120
Epoch 13/20
16/16 [=====] - 1s 58ms/step - loss: 0.1689 - accurac
y: 0.9531 - val_loss: 0.9579 - val_accuracy: 0.8060
Epoch 14/20
16/16 [=====] - 1s 44ms/step - loss: 0.1535 - accurac
y: 0.9536 - val_loss: 0.9896 - val_accuracy: 0.8100
Epoch 15/20
16/16 [=====] - 1s 70ms/step - loss: 0.1403 - accurac
y: 0.9553 - val_loss: 1.0700 - val_accuracy: 0.7920
Epoch 16/20
16/16 [=====] - 1s 43ms/step - loss: 0.1324 - accurac
y: 0.9558 - val_loss: 1.0491 - val_accuracy: 0.8060
Epoch 17/20
16/16 [=====] - 1s 68ms/step - loss: 0.1261 - accurac
y: 0.9584 - val_loss: 1.0587 - val_accuracy: 0.7950
Epoch 18/20
16/16 [=====] - 1s 41ms/step - loss: 0.1196 - accurac
y: 0.9568 - val_loss: 1.0656 - val_accuracy: 0.8060
Epoch 19/20
16/16 [=====] - 1s 59ms/step - loss: 0.1155 - accurac
y: 0.9564 - val_loss: 1.0875 - val_accuracy: 0.8060
Epoch 20/20
16/16 [=====] - 1s 61ms/step - loss: 0.1154 - accurac
y: 0.9575 - val_loss: 1.1097 - val_accuracy: 0.7970
```




## Plot the result

```
In [62]: # Plot the training and validation LOSS
loss = history.history["loss"]
val_loss = history.history["val_loss"]
epochs = range(1, len(loss) + 1)
plt.plot(epochs, loss, "bo", label="Training loss")
plt.plot(epochs, val_loss, "b", label="Validation loss")
plt.title("Training and validation loss")
plt.xlabel("Epochs")
plt.ylabel("Loss")
plt.legend()
plt.show()
```



```
In [63]: # Plot the training and validation ACCURACY
plt.clf()
acc = history.history["accuracy"]
val_acc = history.history["val_accuracy"]
plt.plot(epochs, acc, "bo", label="Training accuracy")
plt.plot(epochs, val_acc, "b", label="Validation accuracy")
plt.title("Training and validation accuracy")
plt.xlabel("Epochs")
plt.ylabel("Accuracy")
plt.legend()
plt.show()
```



## Model Fitting

```
In [64]: # Fit model
model.fit(x_train,
          y_train,
          epochs=9,
          batch_size=512)
```

Epoch 1/9

2023-04-16 23:10:36.980548: W tensorflow/core/common\_runtime/forward\_type\_inference.cc:231] Type inference failed. This indicates an invalid graph that escaped type checking. Error message: INVALID\_ARGUMENT: expected compatible input types, but input 1:

type\_id: TFFT\_OPTIONAL

args {

type\_id: TFFT\_PRODUCT

args {

type\_id: TFFT\_TENSOR

args {

type\_id: TFFT\_BOOL

}

}

}

is neither a subtype nor a supertype of the combined inputs preceding it:

type\_id: TFFT\_OPTIONAL

args {

type\_id: TFFT\_PRODUCT

args {

type\_id: TFFT\_TENSOR

args {

type\_id: TFFT\_LEGACY\_VARIANT

}

}

}

while inferring type of node 'categorical\_crossentropy/cond/output/\_1

0'

```

-----
InvalidArgumentError                                Traceback (most recent call last)
Input In [64], in <cell line: 2>()
      1 # Fit model
----> 2 model.fit(x_train,
      3           y_train,
      4           epochs=9,
      5           batch_size=512)

File ~/opt/anaconda3/lib/python3.9/site-packages/keras/utils/traceback_utils.py:67, in filter_traceback.<locals>.error_handler(*args, **kwargs)
     65 except Exception as e: # pylint: disable=broad-except
     66     filtered_tb = _process_traceback_frames(e.__traceback__)
--> 67     raise e.with_traceback(filtered_tb) from None
     68 finally:
     69     del filtered_tb

File ~/opt/anaconda3/lib/python3.9/site-packages/tensorflow/python/eager/execute.py:54, in quick_execute(op_name, num_outputs, inputs, attrs, ctx, name)
     52 try:
     53     ctx.ensure_initialized()
--> 54     tensors = pywrap_tfe.TFE_Py_Execute(ctx._handle, device_name, op_name,
e,
     55                                         inputs, attrs, num_outputs)
     56 except core._NotOkStatusException as e:
     57     if name is not None:

InvalidArgumentError: Graph execution error:

Detected at node 'gradient_tape/categorical_crossentropy/mul/BroadcastGradientArgs' defined at (most recent call last):
  File "/Users/anjanibonda/opt/anaconda3/lib/python3.9/runpy.py", line 197, in _run_module_as_main
    return _run_code(code, main_globals, None,
  File "/Users/anjanibonda/opt/anaconda3/lib/python3.9/runpy.py", line 87, in _run_code
    exec(code, run_globals)
  File "/Users/anjanibonda/opt/anaconda3/lib/python3.9/site-packages/ipykernel_launcher.py", line 16, in <module>
    app.launch_new_instance()
  File "/Users/anjanibonda/opt/anaconda3/lib/python3.9/site-packages/traitlets/config/application.py", line 846, in launch_instance
    app.start()
  File "/Users/anjanibonda/opt/anaconda3/lib/python3.9/site-packages/ipykernel/kernelapp.py", line 677, in start
    self.io_loop.start()
  File "/Users/anjanibonda/opt/anaconda3/lib/python3.9/site-packages/tornado/platform/asyncio.py", line 199, in start
    self.asyncio_loop.run_forever()
  File "/Users/anjanibonda/opt/anaconda3/lib/python3.9/asyncio/base_events.py", line 596, in run_forever
    self._run_once()
  File "/Users/anjanibonda/opt/anaconda3/lib/python3.9/asyncio/base_events.py", line 1890, in _run_once
    handle._run()
  File "/Users/anjanibonda/opt/anaconda3/lib/python3.9/asyncio/events.py", line 80, in _run
    self._context.run(self._callback, *self._args)
  File "/Users/anjanibonda/opt/anaconda3/lib/python3.9/site-packages/ipykernel/kernelbase.py", line 471, in dispatch_queue

```

```

        await self.process_one()
    File "/Users/anjanibonda/opt/anaconda3/lib/python3.9/site-packages/ipykernel/kernelbase.py", line 460, in process_one
        await dispatch(*args)
    File "/Users/anjanibonda/opt/anaconda3/lib/python3.9/site-packages/ipykernel/kernelbase.py", line 367, in dispatch_shell
        await result
    File "/Users/anjanibonda/opt/anaconda3/lib/python3.9/site-packages/ipykernel/kernelbase.py", line 662, in execute_request
        reply_content = await reply_content
    File "/Users/anjanibonda/opt/anaconda3/lib/python3.9/site-packages/ipykernel/ipkernel.py", line 360, in do_execute
        res = shell.run_cell(code, store_history=store_history, silent=silent)
    File "/Users/anjanibonda/opt/anaconda3/lib/python3.9/site-packages/ipykernel/zmqshell.py", line 532, in run_cell
        return super().run_cell(*args, **kwargs)
    File "/Users/anjanibonda/opt/anaconda3/lib/python3.9/site-packages/IPython/core/interactiveshell.py", line 2880, in run_cell
        result = self._run_cell(
    File "/Users/anjanibonda/opt/anaconda3/lib/python3.9/site-packages/IPython/core/interactiveshell.py", line 2935, in _run_cell
        return runner(coro)
    File "/Users/anjanibonda/opt/anaconda3/lib/python3.9/site-packages/IPython/core/async_helpers.py", line 129, in _pseudo_sync_runner
        coro.send(None)
    File "/Users/anjanibonda/opt/anaconda3/lib/python3.9/site-packages/IPython/core/interactiveshell.py", line 3134, in run_cell_async
        has_raised = await self.run_ast_nodes(code_ast.body, cell_name,
    File "/Users/anjanibonda/opt/anaconda3/lib/python3.9/site-packages/IPython/core/interactiveshell.py", line 3337, in run_ast_nodes
        if await self.run_code(code, result, async_=asy):
    File "/Users/anjanibonda/opt/anaconda3/lib/python3.9/site-packages/IPython/core/interactiveshell.py", line 3397, in run_code
        exec(code_obj, self.user_global_ns, self.user_ns)
    File "/var/folders/0c/spq36xkd5vz7k9940sd5prhw0000gn/T/ipykernel_2957/1659230057.py", line 2, in <cell line: 2>
        model.fit(x_train,
    File "/Users/anjanibonda/opt/anaconda3/lib/python3.9/site-packages/keras/utils/traceback_utils.py", line 64, in error_handler
        return fn(*args, **kwargs)
    File "/Users/anjanibonda/opt/anaconda3/lib/python3.9/site-packages/keras/engine/training.py", line 1409, in fit
        tmp_logs = self.train_function(iterator)
    File "/Users/anjanibonda/opt/anaconda3/lib/python3.9/site-packages/keras/engine/training.py", line 1051, in train_function
        return step_function(self, iterator)
    File "/Users/anjanibonda/opt/anaconda3/lib/python3.9/site-packages/keras/engine/training.py", line 1040, in step_function
        outputs = model.distribute_strategy.run(run_step, args=(data,))
    File "/Users/anjanibonda/opt/anaconda3/lib/python3.9/site-packages/keras/engine/training.py", line 1030, in run_step
        outputs = model.train_step(data)
    File "/Users/anjanibonda/opt/anaconda3/lib/python3.9/site-packages/keras/engine/training.py", line 893, in train_step
        self.optimizer.minimize(loss, self.trainable_variables, tape=tape)
    File "/Users/anjanibonda/opt/anaconda3/lib/python3.9/site-packages/keras/optimizers/optimizer_v2/optimizer_v2.py", line 537, in minimize
        grads_and_vars = self._compute_gradients(
    File "/Users/anjanibonda/opt/anaconda3/lib/python3.9/site-packages/keras/optimizers/optimizer_v2/optimizer_v2.py", line 590, in _compute_gradients

```

```

grads_and_vars = self._get_gradients(tape, loss, var_list, grad_loss)
File "/Users/anjanibonda/opt/anaconda3/lib/python3.9/site-packages/keras/opt
timizers/optimizer_v2/optimizer_v2.py", line 471, in _get_gradients
grads = tape.gradient(loss, var_list, grad_loss)
Node: 'gradient_tape/categorical_crossentropy/mul/BroadcastGradientArgs'
Incompatible shapes: [512] vs. [512,46]
[[{{node gradient_tape/categorical_crossentropy/mul/BroadcastGradient
Args}}]] [Op:__inference_train_function_6274]

```

```

In [35]: # Evaluate model
results = model.evaluate(x_test, y_test)

```

```

71/71 [=====] - 0s 2ms/step - loss: 1.3359 - accurac
y: 0.7876

```

```

In [37]: # Show results
print(f'Test loss: {results[0]:0.3f}\nTest accuracy: {results[1]:0.3f}')

```

```

Test loss: 1.336
Test accuracy: 0.788

```

```

In [38]: # Predict test data
predictions = model.predict(x_test)

```

```

In [39]: predictions[0]

```

```

Out[39]: array([4.96593930e-06, 8.68394636e-07, 1.36702042e-07, 8.62344682e-01,
 7.67317712e-02, 1.47404214e-10, 3.00806868e-09, 6.34642565e-05,
 6.04539830e-03, 1.01301430e-06, 4.23747331e-08, 2.77778786e-02,
 3.22184064e-06, 5.35935169e-06, 3.60176813e-08, 2.30063506e-07,
 6.53075171e-04, 4.67550336e-08, 1.18455284e-06, 2.89977156e-03,
 2.32233517e-02, 1.07884094e-04, 2.67851181e-08, 6.77256153e-08,
 4.97829822e-09, 4.14858590e-07, 2.04205611e-10, 1.11060615e-07,
 4.64958021e-06, 3.71215856e-06, 1.65704259e-05, 1.66338974e-11,
 9.97571419e-07, 9.28010724e-10, 5.14841759e-06, 1.58880198e-07,
 7.74415894e-05, 4.70382240e-08, 4.13131920e-06, 2.11911884e-05,
 3.13704675e-07, 5.16779664e-07, 7.80447973e-09, 1.84905355e-10,
 2.11908702e-12, 1.00457470e-07], dtype=float32)

```

## Assignment 5.3

Implement the housing price regression model found in section 3.6 of Deep Learning with Python

### Load Data

```

In [50]: # Load the Boston Housing dataset
from keras.datasets import boston_housing
# split train / test data
(train_data, train_targets), (test_data, test_targets) = boston_housing.load_data

```

```

Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datas
ets/boston_housing.npz
57026/57026 [=====] - 0s 1us/step

```

```

In [51]: print("The length of training dataset: {}".format(len(train_data)))
print("The length of test dataset: {}".format(len(test_data)))

```

The length of training dataset: 404  
The length of test dataset: 102

## Data Preparation

```
In [52]: ## Perform feature-wise normalization  
# Normalize train data  
mean = train_data.mean(axis=0)  
train_data -= mean  
std = train_data.std(axis=0)  
train_data /= std  
# Normalize test data  
test_data -= mean  
test_data /= std
```

## Model Building

```
In [53]: def build_model():  
    # Define the model  
    model = models.Sequential()  
    model.add(layers.Dense(64, activation='relu',  
                           input_shape=(train_data.shape[1],)))  
    model.add(layers.Dense(64, activation='relu'))  
    model.add(layers.Dense(1))  
    # Compile the model  
    model.compile(optimizer='rmsprop', loss='mse', metrics=['mae'])  
    return model
```

## Validate the model

```
In [54]: # Set up K-fold validation  
k = 4  
num_val_samples = len(train_data) // k  
num_epochs = 500  
all_mae_histories = []  
for i in range(k):  
    print(f"Processing fold #: {i}")  
    # prepare the validation data (from partition #k)  
    val_data = train_data[i * num_val_samples: (i + 1) * num_val_samples]  
    val_targets = train_targets[i * num_val_samples: (i + 1) * num_val_samples]  
  
    # prepare training data (from non-k partitions)  
    partial_train_data = np.concatenate(  
        [train_data[:i * num_val_samples],  
         train_data[(i + 1) * num_val_samples:]],  
        axis=0)  
    partial_train_targets = np.concatenate(  
        [train_targets[:i * num_val_samples],  
         train_targets[(i + 1) * num_val_samples:]],  
        axis=0)  
  
    # build the already compiled keras model  
    model = build_model()  
  
    # train the model, saving validation logs at each fold  
    history = model.fit(partial_train_data, partial_train_targets,  
                        validation_data=(val_data, val_targets),
```

```

epochs=num_epochs, batch_size=16, verbose=0) # verbose=
mae_history = history.history["val_mae"]
all_mae_histories.append(mae_history)

```

```

Processing fold #: 0
Processing fold #: 1
Processing fold #: 2
Processing fold #: 3

```

```

In [55]: # Compute the average of per-epoch MAE scores for all folds
average_mae_history = [
    np.mean([x[i] for x in all_mae_histories]) for i in range(num_epochs)
]


```

## Plot the results

```

In [56]: # Plot the validation MAE by epoch
plt.plot(range(1, len(average_mae_history) + 1), average_mae_history)
plt.xlabel("Epochs")
plt.ylabel("Validation MAE")
plt.show()


```



```

In [57]: # Plot again, omitting the first 10 data points
truncated_mae_history = average_mae_history[10:]
plt.plot(range(1, len(truncated_mae_history) + 1), truncated_mae_history)
plt.xlabel("Epochs")
plt.ylabel("Validation MAE")
plt.show()

```



## Model Training

```
In [58]: model = build_model() # get a fresh compiled model
# Train on the entirety of the training data
model.fit(train_data, train_targets,
          epochs=130, batch_size=16, verbose=0)
test_mse_score, test_mae_score = model.evaluate(test_data, test_targets)

4/4 [=====] - 0s 3ms/step - loss: 18.8116 - mae: 2.5601
```

```
In [59]: # Evaluate model
test_mae_score
```

```
Out[59]: 2.5600898265838623
```

## Generate Predictions

```
In [60]: # Predict test data
predictions = model.predict(test_data)
predictions[0]

4/4 [=====] - 0s 3ms/step
Out[60]: array([8.194735], dtype=float32)
```

```
In [ ]:
```